

# MNIST

April 11, 2023

```
[ ]: import tensorflow as tf
import tensorflow_datasets as tfds
import numpy as np
import matplotlib.pyplot as plt
```

2023-04-10 19:16:27.224801: I tensorflow/core/platform/cpu\_feature\_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: SSE3 SSE4.1 SSE4.2 AVX AVX2 FMA To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

```
[ ]: (ds_train, ds_test), ds_info = tfds.load(
    'mnist',
    split=['train', 'test'],
    shuffle_files=True,
    as_supervised=True,
    with_info=True,
)
```

```
[ ]: def normalize_img(image, label):
    """Normalizes images: `uint8` -> `float32`."""
    return tf.cast(image, tf.float32) / 255., label

ds_train = ds_train.map(
    normalize_img, num_parallel_calls=tf.data.AUTOTUNE)
ds_train = ds_train.cache()
ds_train = ds_train.shuffle(ds_info.splits['train'].num_examples)
ds_train = ds_train.batch(128)
ds_train = ds_train.prefetch(tf.data.AUTOTUNE)
```

```
[ ]: ds_test = ds_test.map(
    normalize_img, num_parallel_calls=tf.data.AUTOTUNE)
ds_test = ds_test.batch(128)
ds_test = ds_test.cache()
ds_test = ds_test.prefetch(tf.data.AUTOTUNE)
```

```
[ ]: model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    #tf.keras.layers.GlobalAveragePooling2D(),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10)
])
model.compile(
    optimizer=tf.keras.optimizers.Adam(0.001),
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    metrics=[tf.keras.metrics.SparseCategoricalAccuracy(name='accuracy')],
)

testing = model.fit(
    ds_train,
    epochs=10,
    validation_data=ds_test,
)
```

Epoch 1/10

469/469 [=====] - 5s 4ms/step - loss: 0.2700 - accuracy: 0.9196 - val\_loss: 0.1287 - val\_accuracy: 0.9610

Epoch 2/10

469/469 [=====] - 2s 3ms/step - loss: 0.1117 - accuracy: 0.9657 - val\_loss: 0.1154 - val\_accuracy: 0.9652

Epoch 3/10

469/469 [=====] - 2s 3ms/step - loss: 0.0760 - accuracy: 0.9775 - val\_loss: 0.0986 - val\_accuracy: 0.9705

Epoch 4/10

469/469 [=====] - 2s 3ms/step - loss: 0.0553 - accuracy: 0.9828 - val\_loss: 0.0977 - val\_accuracy: 0.9720

Epoch 5/10

469/469 [=====] - 2s 3ms/step - loss: 0.0419 - accuracy: 0.9870 - val\_loss: 0.0980 - val\_accuracy: 0.9733

Epoch 6/10

469/469 [=====] - 2s 3ms/step - loss: 0.0336 - accuracy: 0.9893 - val\_loss: 0.0924 - val\_accuracy: 0.9753

Epoch 7/10

469/469 [=====] - 2s 3ms/step - loss: 0.0264 - accuracy: 0.9920 - val\_loss: 0.1005 - val\_accuracy: 0.9734

Epoch 8/10

469/469 [=====] - 2s 3ms/step - loss: 0.0215 - accuracy: 0.9930 - val\_loss: 0.1006 - val\_accuracy: 0.9752


Epoch 9/10

469/469 [=====] - 2s 3ms/step - loss: 0.0188 - accuracy: 0.9940 - val\_loss: 0.0993 - val\_accuracy: 0.9768

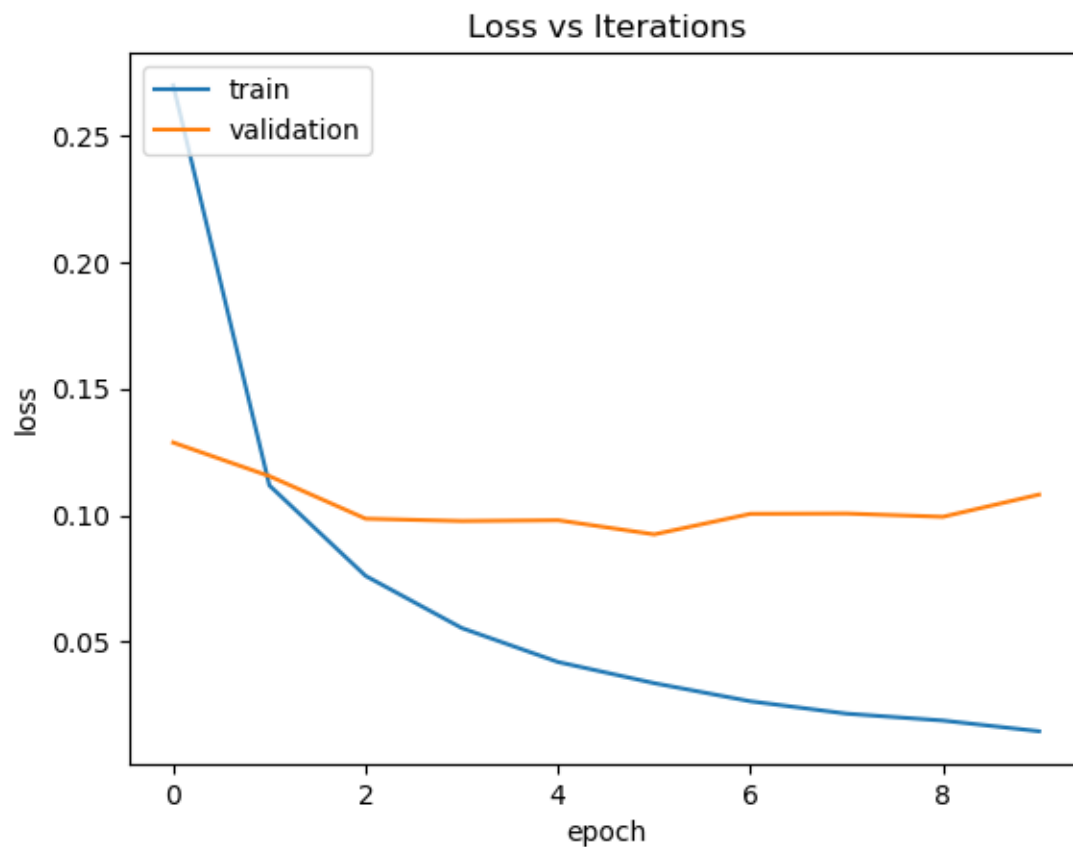
Epoch 10/10

469/469 [=====] - 1s 3ms/step - loss: 0.0145 -  
accuracy: 0.9956 - val\_loss: 0.1081 - val\_accuracy: 0.9739

```
[ ]: plt.plot(testing.history['accuracy'])  
plt.plot(testing.history['val_accuracy'])  
plt.title('Accuracy vs Iterations')  
plt.ylabel('accuracy')  
plt.xlabel('epoch')  
plt.legend(['train', 'validation'], loc='upper left')  
plt.show()
```



```
[ ]: plt.plot(testing.history['loss'])  
plt.plot(testing.history['val_loss'])  
plt.title('Loss vs Iterations')  
plt.ylabel('loss')  
plt.xlabel('epoch')  
plt.legend(['train', 'validation'], loc='upper left')  
plt.show()
```



```
[ ]: model.evaluate(ds_test)
```

```
79/79 [=====] - 0s 2ms/step - loss: 0.1081 - accuracy: 0.9739
```

```
[ ]: [0.10812117904424667, 0.9739000201225281]
```